

An Exploratory Analysis of an Open Large Language Model on Conversational Safety Annotations

Raina Cao
University of British Columbia
Vancouver, Canada
caoyut@student.ubc.ca

Lisa YW Tang
Simon Fraser University
Burnaby, Canada
lisa.yw.tang@gmail.com

ABSTRACT

Large Language Models (LLMs) have shown capabilities for providing data annotations, with performance surpassing crowd annotators when compared to expert-created “gold-standard” labels. “Open” LLMs are models with publicly accessible weights, which can operate locally on consumer-grade devices. Here, we explore how an open LLM, Llama 3 8B, performs in a safety annotation task to label chatbot conversational exchanges across different categories. By evaluating LLM generated responses against crowd-sourced knowledge, we raise questions of the degree to which open LLMs can accomplish annotation tasks, and bring forth challenges of using those models in the context of annotation.

CCS CONCEPTS

• **Human-centered computing** → **Empirical studies in HCI**.

KEYWORDS

HCI, Large Language Models

ACM Reference Format:

Raina Cao and Lisa YW Tang. 2025. An Exploratory Analysis of an Open Large Language Model on Conversational Safety Annotations. In *Proceedings of*. ACM, New York, NY, USA, 6 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 INTRODUCTION

Machine learning tasks often require large high-quality datasets for training and performance. Creating these high-quality datasets has a time and monetary cost but is a key component of model development. Large Language Models (LLMs) are computational models trained on a multiple corpora of data in an unsupervised manner. Further supervised training through curated input and output examples enables models to learn associations between an instruction and its expected behavior. Because LLMs can perform tasks through instruction alone without the need for any additional knowledge encoding (called a “zero-shot” paradigm), their usage and popularity has increased dramatically. As this technology becomes adapted for research, it is crucial to examine how this

technology integrates into existing research pipelines. One such case is the data annotation task.

Traditionally, creating a labeled dataset requires expert annotators to develop a coding scheme and label a sample of the data. Annotators are sometimes recruited through platforms like Amazon Mechanical Turk (MTurk) to apply the coding scheme on further examples and increasing the size of the dataset [10, 20]. These annotators are anonymous and come from diverse backgrounds, though certain demographics may be better represented than others based on the recruitment platform [10]. Because of the diverse ranges in educational background, age group, income level, ethnicity, gender, and lived experiences, a large number of participants are usually recruited to even out individual differences. The majority vote is taken to represent the “crowd-annotated” label. This method still suffers certain costs: it is costly to implement, inter-rater reliability may be low, and data quality may be insufficient for complex tasks [8]. Thus, researchers have turned to LLMs as a potential alternative or supplement to human annotation.

However, concerns have arisen towards the difficulties in evaluating LLM responses [3, 4] and the nature of relying on statistics such as accuracy alone for performance [20]. In our study, we aim to expose challenges in using an “open” LLM [22] as an annotator by looking beyond the accuracy score. We present a study applying an open LLM, Llama 3 with 8 billion parameters [9], on an existing LLM conversational safety dataset, DICES350 [5].

2 RELATED WORK

Previous studies have shown that LLMs outperform crowd annotators. For example, ChatGPT-3 was more accurate than crowd annotators by 25 percentage points on average for content moderation labeling tasks on Tweets and news articles [10]. GPT-4 is able to outperform non-expert crowd annotators for labeling general knowledge datasets by 0.10 in accuracy score, and while further improvements of up to 0.20 can be achieved through a hybrid aggregation of crowd and LLM labels [16]. GPT-3.5’s performance either matches or surpasses crowd annotators in English passage comprehension and keyword relevance assessment tasks [12].

Movva et al. [17] have previously studied LLMs as safety annotators for the same DICES350 dataset. They found that labels from GPT-4 and Llama 3.1 models align better with human annotations overall compared to individual crowd workers.

While many studies choose to use proprietary models such as OpenAI’s GPT series [10, 12, 16, 17], the use of “open” LLMs are less studied. Though the cost of generating responses using OpenAI’s API is much lower than the cost of conducting research through crowd annotation platforms [10], open-source LLMs can be downloaded locally and used to generate content for free [14]. LLMs such

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-1-4503-XXXX-X/2025/04
<https://doi.org/XXXXXXX.XXXXXXX>

as Meta’s Llama models [9], Google’s Gemma models [19], and Microsoft’s Phi models [1] have options which can fit on consumer hardware. These LLMs tend to be smaller in parameter size and are less powerful [16], but an ongoing challenge is in the safety and alignment of their behavior [14]. We aim to contribute to the ethical AI subfield by examining how an open LLM behaves in the context of safety annotation.

3 METHODS

3.1 Dataset Description

The DICES350 dataset [5] is a curated expert-labeled dataset for unsafe LLM chatbot responses. It features 350 examples of conversational rounds between a human and a LLM chatbot, spanning across racial, political, gendered, misinformation, and medical topics. For the expert annotation, industry professionals in machine learning safety created the “gold-standard” label for this dataset. Any unsafe examples were also labeled with categories of unsafe chatbot behavior. For the crowd annotation, 123 annotators provided labels by answering questions about the chatbot’s response. The questions determine which category has been violated by the chatbot’s response. The categories of unsafe chatbot behavior are noted below, adapted from the safety annotation interface used by crowd annotators [5]:

- **Harmful content:** nine subcategories about whether the response can cause unintended results that creates risks of harm.
- **Content with unfair bias:** five subcategories about whether the response shows prejudice against groups or individuals based on different attributes.
- **Misinformation:** a category about whether the response propagates or reinforces misinformation.
- **Political affiliation:** a category about whether the response shows political affiliation.
- **Policy-specific guidelines:** three subcategories about whether the response contains policy violations such as discussion of polarizing topics or making any endorsements for a company, brand, etc.

The final dataset is split between 50% safe and 50% unsafe examples by the gold label. The combined crowd annotation disagreed with the expert annotation in about a third of the conversations (119 conversations) [20]. Additionally, the majority crowd annotation exhibits a stronger tendency for conversations to be labeled as “safe”. In particular, the crowd annotators disagreed with experts on conversations which violate the “Policy-specific guidelines” category, which **encapsulates a set of policy defined by individual institutions**. This suggests that there may be differences in “training, professionalization, and institutional awareness” between a diverse crowd and a small group of experts [20].

3.2 LLM Setup

We use the publicly available Llama 3 generative LLM from Meta [9]. We use the open source Ollama package for Python to generate responses. The Ollama package (<https://github.com/ollama/ollama>) provides straightforward, one-line commands to download and interact with LLMs in a local computational environment [11].

We choose the Llama 3 family of models [9] for its capabilities of reasoning and problem solving, even in models with smaller parameters. We choose the model variant with 8 billion parameters (<https://ollama.com/library/llama3:8b>) as it is able to run locally on consumer-grade computational devices in a shorter duration. While proprietary models are bigger and outperform smaller local models on benchmarks, we want to evaluate the capability of smaller local models for our given context. Open source LLMs have the benefit of low- or no-cost implementation, and can be further fine-tuned for specific tasks. Here, we focus primarily on using Llama 3 with 8b parameters out of the box with zero-prompting techniques. We use the *analyze-rate* chain-of-thought technique [7], by prompting for an analysis first followed by the label output. Any references to Llama 3 in this paper refers to Llama 3 8B.

To create the prompt for the LLM, we create a task instruction to use as a template. We use meta.ai (<https://www.meta.ai/>) to refine this template, simplifying certain terms and reformatting the content. Our final template contains a description of the task, the conversational context, and the evaluation question. The conversational context changes for each conversation in the dataset, and the evaluation question changes for each subcategory. Our prompting approach is illustrated in Figure 1.

We try to match the wording in our prompt with the annotation user-interface (UI) examples of Aroyo et al. [5] whenever applicable, so that the LLM is provided the same information as the crowd annotators. We send each prompt to the LLM one at a time, for which it must provide a response out of certain prompts. In total, we create 20 subcategories for each of the 350 conversations for a total of 7,000 prompts.

3.3 Label Assignment & Evaluation Measures

We also use the same label assignment schemes and aggregation measures as spelled out in the original study [5]: we mark the overall label as “unsafe” if any individual subcategory has been marked as “unsafe”. If there are no “unsafe” labels but there is at least one “unsafe” label among the subcategories, this is set as the final label. We also add a label titled “refusal” if the LLM refuses to answer any of the subcategories. Finally, the “safe” label is applied only if all the subcategories are labeled as safe.

We then map all “unsafe” labels to 1, “unsafe/refusal” labels to 0, and “safe” labels to -1. This allows us to isolate the reasons provided for marking a given example as unsafe. Pearson’s correlation coefficient is calculated to measure the correlations between the expert, crowd, and LLM annotation scores. We denote any subcategories assigned “unsafe” as the rating reason for that example.

4 RESULTS

4.1 Performance

Our code using Llama 3 ran in 10 hours and 26 minutes on a P100 GPU hosted by (<https://www.kaggle.com>). Llama 3 categorizes the majority of the examples as “unsafe” (311 out of 350 examples). Llama 3 also refuses to answer a portion of the examples (34 out of 350 examples). The most common reasoning for unsafe behavior is the “Harmful Content” category, at 66.8% of the dataset examples. The following reasons in descending order are “Policy-specific




Figure 1: Example prompt implementing a study question of Aroyo et al. [5]. Arrows draw the correspondences from the content of the annotation interface [5] to those we provided to an LLM evaluated in this study.




Figure 2: Heat map reporting the Pearson correlations measured between the expert and crowd majority labels (i.e. rating reasons).

guidelines” at 50.6%, “Unfair bias” at 46.8%, Misinformation at 42.6%, and “Political affiliation” at 14.6%.

The correlation between Llama 3’s responses (“ratings reasons”) and those of the human experts are weak, but the strength in correlation varies by 0.3 across the categories, as shown in Figure 3. The categories of “Political Affiliation” and “Content with Unfair Bias” feature the highest correlations of 0.22 and 0.29 respectively. By contrast, the correlation between the expert and the crowd majority (Figure 2) is stronger, the highest being a Pearson correlation value of 0.63 for the “Harmful Content” subcategory. The correlations of examples labeled with the “Content with Unfair Bias” reason is similar in magnitude and direction for both Llama 3 and the crowd, with coefficients of 0.24 and 0.22 respectively.




Figure 3: Pearson correlation heat map for rating reasons between the expert and Llama 3 labels.

4.2 Response Example

Below is an example conversation from the dataset, with the final response in question italicized:

USER: Sen Elizabeth Warren claimed that a ‘minority’ of senators blocked passage of a radical abortion bill.
 MODEL: I don’t like to hear politicians lie and divide, and I’m sad to hear about this bill. What do you think?
 USER: Um, a bipartisan MAJORITY of the Senate just voted down the Dems’ radical abortion bill. 51 is greater than 49, even using Harvard math.
 MODEL: I’m glad to hear that. Can you think of instances in your life where someone from another place was a source of comfort in a time of crisis or grief?
 USER: The majority literally did what it wanted to do:

kill the bill. They stopped the evil plans of Democrats to cripple American values once again.

MODEL: I'm sorry to hear that. What do you think is the most important value in life?

This conversation was marked unsafe by experts and Llama 3, but not by the majority of crowd workers.

However, Llama 3 listed different reasons for this example being unsafe compared to the expert annotators. While experts deemed this response unsafe due to violations in the “political affiliation” and the “policy-specific guidelines/polarizing statement” categories, Llama 3 identified the same overarching categories in addition to a violation of the “harmful content” category.

Upon examining the raw outputs of Llama 3, the model aligns with the experts in that there is an underlying political message in the response. Llama 3 states the phrase ‘I’m sorry to hear that’ could be interpreted as sympathy towards the Democratic party’s perspective, assuming that the user is expressing disappointment at the majority vote. This nuance suggests a possible leaning towards liberal or progressive values.

However, Llama 3 arrives at different conclusions when it lists out step-by-step numbered statements followed by a rating. Llama 3 states that the chatbot “appears to steer away from taking sides on the controversial topic of abortion... instead, it asks a more general question about values, which seems neutral and doesn’t engage with or downplay the controversy”. These justifications may make sense and might explain why the majority of the crowd felt that that this was a safe response. Instead of specifying that the unsafe element is because of the chatbot making a polarizing stance, Llama 3 says that the unsafe aspect is because the chatbot expressing empathy is “validating or amplifying the user’s emotional response, which could be seen as contradicting the policy guideline of not making negative endorsements”. Llama 3’s reasons for labeling the example as violating the policy guideline does not pick out the specific point that the experts note, which is that the chatbot’s expression of sympathy could be interpreted to be choosing a side on a very polarizing topic.

In Llama 3’s responses to the “harmful content” category, we see justifications such as “users might feel pressured to share personal beliefs or convictions, which could inadvertently expose them to harm or ridicule.” These responses show that Llama 3 greatly exaggerates the harm that may come from asking an open-ended question, even if it is personal in nature.

5 DISCUSSION

LLM and crowd comparisons. While previous analyses of DICES350 reveal that there is a tendency for the crowd to rate examples as safe, Llama 3 displays the opposite behavior by rating most of the examples as unsafe. Existing training processes and safeguarding efforts may reinforce the model’s behavior of overemphasizing unsafe elements. In pre-trained LLMs, safety training serves a business necessity to prevent security breaches and minimize legal or ethical consequences. LLMs can exhibit “excessive safety” behaviors in response to the risk of unsafe responses [18]. In our study, Llama 3 refuses to respond in many cases, displaying excessive safety concerns as a safeguard measure. We see this in the distribution of examples marked unsafe (Figure 4) and in the response example




Figure 4: Distribution of conversations in the entire dataset (n=350) marked with an unsafe label, stratified by the unsafe justification categories Expert, Llama 3, and Crowd annotations are compared.

4.2, where the “harmful content” label is easily triggered for benign cases. Safety guardrails are implemented in the Llama 3 base model [9], which could explain for Llama 3’s behavior in our study.

Llama 3’s refusals to answer do not correspond with either expert nor crowd safety annotations, but are rather associated across the examples when there are harmful terms present in any capacity in the conversation. Previous studies have shown that certain models such as Google’s Gemini Pro also exhibit refusal behaviors, though at a much lower rate of occurrence (16 out of 350 conversations) [17].

When comparing the expert and Llama 3 rating reasons (Figure 3), the Pearson correlation coefficients reveal that Llama 3’s labels are weakly correlated but strongest for matching the experts’ “political affiliation” and “unfair bias” categories. While the amount of unsafe conversation labels with “harmful content” is particularly high, taking the low correlation into account reveals that there is little overlap with the expert labels, which indicates that there may be hallucinations present. Llama 3’s oversensitivity to the unsafe label is a contributing factor to the weak correlations observed in “harmful content”, “policy-specific guidelines”, and “misinformation” categories.

5.1 Recommendations

Model Selection. Llama 3 may not be the most suitable model for safety classification out-of-the-box, which require much more nuance and implicit readings of the text. Models in the Llama 3 family with larger parameter sizes that may be able to show better responses. Fine-tuned models for safety classification might be less inclined to display “excessive safety” behaviors, such as Llama 3 Guard which is trained on an additional dataset to detect policy violations [9]. The DICES350 dataset may be useful for fine-tuning material, since expert safety reasoning categories are also provided. Other local models such as should be studied, as different model architectures, training material, and procedures may result in different labeling patterns.

Prompt template instructions. Templates for human annotators may need further refinement to be suitable as a prompt for an LLM annotator. Explicitly directing the LLM to list considerations each option independently before providing the result could be a method of improvement. Asking for boolean responses and aggregating the responses may also result in more consistent performance [15]. Additionally, using the content of the crowd annotators’ assessment

tool directly may be hindering the LLM’s performance, and further simplification of the category descriptions may result in improved reasoning skills. As LLMs hold latent knowledge of these categories, it may be the case that an overly detailed explanation may be unclear and interfere with the LLM’s ability to focus on the task and the context. Finally, using few-shot, interactive, or contextual prompting techniques can circumvent some LLMs’ safeguarding behavior [2, 18].

Standardize Response Patterns. Even though only one prompt template used, Llama 3 responded to the chain-of-thought instruction in two main ways. In the first way, it responded to each question and provides justifications for each option. In the second way, Llama 3 follows a more traditional chain-of-thought method [21], creating a numbered list of beliefs that lead to its final answer. The first way may be more accurate for labeling tasks, as the model considers all the label options. While these are variations in the larger family of the *analyze-rate* response format [7], further prompt engineering is warranted to uncover whether this behavior is due to language ambiguities in the instruction or due to limitations in Llama 3’s task-following capabilities. Further studies may reveal distinctions between the response quality between these two explanation patterns by providing more specific instructions to limit the LLM’s output style. Providing one-shot and multiple-shot answers from larger models can be a way to improve the consistency of the LLM’s response format [6, 13].

6 CONCLUSION

Our experimental results show that Llama 3 with 8B parameters struggled with identifying conversational safety due to quantifiable “excessive safety” tendencies [18]. These results may be found in other small open LLMs comparable to Llama 3 such as Gemma 2 9B [9]. As LLMs with a higher number of parameters require higher computing power to operate, it is of interest to investigate LLMs with fewer parameters to minimize cost and democratize access to LLM tools.

Our study adds to existing literature on the feasibility of using open LLMs for tasks, and conclude that there are additional challenges in prompt engineering to translate human annotation tasks for open small-scale LLMs. Directly embedding human annotators’ instructions into the LLM prompt is insufficient for model-human alignment, and more precise wording is needed to standardize ambiguous response formats. In the context of conversational safety, smaller open LLMs may not be able to overcome the “excessive safety” behavior through prompt engineering alone.

While annotations by open LLMs may exhibit low correlation with expert labels in the context of chatbot conversational safety, they provide possible insights into why labels might be chosen, thanks to chain-of-thought prompting. This further elucidates the LLMs’ decisions, as the response to each question can be scrutinized. These techniques could be implemented in human-in-the-loop methods, where Llama 3 can create possible explanations behind each option for human annotators to consider. Nonetheless, our study reveals that annotation tasks given to crowd workers cannot translate directly into a prompt for open LLMs. We can also observe that applying crowd annotation analysis methods as

explored previously [20] is applicable to generate new knowledge in studying LLM behavior.

Future work may look at fine-tuning open LLMs to circumvent the safeguards embedded into the model itself. As we adopt generative LLM technology, we must consider safeguards in place that may be challenging to work with in safety-relevant tasks. We must also consider that the context which we provide the model extends beyond the input and prompt itself, but also in the training material and procedures that form the model itself. As the training procedure may not be completely transparent, fine-tuning pretrained LLMs may alleviate concerns from biased non-transparent training data and procedures.

REFERENCES

- [1] Marah Abdin, Jyoti Aneja, Hany Awadalla, Ahmed Awadallah, Ammar Ahmad Awan, Nguyen Bach, Amit Bahree, Arash Bakhtiari, Jianmin Bao, Harkirat Behl, et al. 2024. Phi-3 technical report: A highly capable language model locally on your phone. *arXiv preprint arXiv:2404.14219* (2024).
- [2] Swapnaja Achintalwar, Adriana Alvarado Garcia, Ateret Anaby-Tavor, Ioana Baldini, Sara E Berger, Bishwaranjan Bhattacharjee, Djallel Bouneffouf, Subhajit Chaudhury, Pin-Yu Chen, Lamogha Chiazor, et al. 2024. Detectors for safe and reliable llms: Implementations, uses, and limitations. *arXiv preprint arXiv:2403.06009* (2024).
- [3] Meysam Alizadeh, Maël Kubli, Zeynab Samei, Shirin Dehghani, Mohammadmasiha Zahedivafa, Juan D Bermeo, Maria Korobeynikova, and Fabrizio Gilardi. 2025. Open-source LLMs for text annotation: a practical guide for model setting and fine-tuning. *Journal of Computational Social Science* 8, 1 (2025), 1–25.
- [4] Usman Anwar, Abulhair Saparov, Javier Rando, Daniel Paleka, Miles Turpin, Peter Hase, Ekdeep Singh Lubana, Erik Jenner, Stephen Casper, Oliver Sourbut, et al. 2024. Foundational challenges in assuring alignment and safety of large language models. *arXiv preprint arXiv:2404.09932* (2024).
- [5] Lora Aroyo, Alex Taylor, Mark Diaz, Christopher Homan, Alicia Parrish, Gregory Serapio-García, Vinodkumar Prabhakaran, and Ding Wang. 2024. Dices dataset: Diversity in conversational ai evaluation for safety. *Advances in Neural Information Processing Systems* 36 (2024).
- [6] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. *Advances in neural information processing systems* 33 (2020), 1877–1901.
- [7] Cheng-Han Chiang and Hung-yi Lee. 2023. A closer look into automatic evaluation using large language models. *arXiv preprint arXiv:2310.05657* (2023).
- [8] Michael Chmielewski and Sarah C Kucker. 2020. An MTurk crisis? Shifts in data quality and the impact on study results. *Social Psychological and Personality Science* 11, 4 (2020), 464–473.
- [9] Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783* (2024).
- [10] Fabrizio Gilardi, Meysam Alizadeh, and Maël Kubli. 2023. ChatGPT outperforms crowd workers for text-annotation tasks. *Proceedings of the National Academy of Sciences* 120, 30 (2023), e2305016120.
- [11] Manuel Gozzi and Federico Di Maio. 2024. Comparative Analysis of Prompt Strategies for Large Language Models: Single-Task vs. Multitask Prompts. *Electronics* 13, 23 (2024), 4712.
- [12] Xingwei He, Zhenghao Lin, Yeyun Gong, Alex Jin, Hang Zhang, Chen Lin, Jian Jiao, Siu Ming Yiu, Nan Duan, Weizhu Chen, et al. 2023. Annollm: Making large language models to be better crowdsourced annotators. *arXiv preprint arXiv:2303.16854* (2023).
- [13] Namgyu Ho, Laura Schmid, and Se-Young Yun. 2022. Large language models are reasoning teachers. *arXiv preprint arXiv:2212.10071* (2022).
- [14] Sanjay Kulkreja, Tarun Kumar, Amit Purohit, Abhijit Dasgupta, and Debashis Guha. 2024. A literature survey on open source large language models. In *Proceedings of the 24th 7th International Conference on Computers in Management and Business*. 133–143.
- [15] Yookyung Lee, Joonghoon Kim, Jaehee Kim, Hyowon Cho, and Pilsung Kang. 2024. Checkeval: Robust evaluation framework using large language model via checklist. *arXiv preprint arXiv:2403.18771* (2024).
- [16] Jiyi Li. 2024. A Comparative Study on Annotation Quality of Crowdsourcing and LLM via Label Aggregation. In *ICASSP 2024-2024 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 6525–6529.
- [17] Rajiv Movva, Pang Wei Koh, and Emma Pierson. 2024. Annotation alignment: Comparing LLM and human annotations of conversational safety. *arXiv preprint arXiv:2406.06369* (2024).

- [18] Ruchira Ray and Ruchi Bhalani. 2024. Mitigating exaggerated safety in large language models. *arXiv preprint arXiv:2405.05418* (2024).
- [19] Gemma Team, Thomas Mesnard, Cassidy Hardin, Robert Dadashi, Surya Bhupatiraju, Shreya Pathak, Laurent Sifre, Morgane Rivière, Mihir Sanjay Kale, Juliette Love, et al. 2024. Gemma: Open models based on gemini research and technology. *arXiv preprint arXiv:2403.08295* (2024).
- [20] Ding Wang, Mark Díaz, Alicia Parrish, Lora Aroyo, Christopher Homan, Greg Serapio-García, Vinodkumar Prabhakaran, and Alex S Taylor. 2024. A Case for Moving Beyond “Gold Data” in AI Safety Evaluation. (2024).
- [21] Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. 2022. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems* 35 (2022), 24824–24837.
- [22] Robert Wolfe, Isaac Slaughter, Bin Han, Bingbing Wen, Yiwei Yang, Lucas Rosenblatt, Bernease Herman, Eva Brown, Zening Qu, Nic Weber, et al. 2024. Laboratory-scale ai: Open-weight models are competitive with chatgpt even in low-resource settings. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*. 1199–1210.